

## Set 5 - MPI Part II

Issued: November 27, 2020

Hand in (optional): December 11, 2020 08:00am


**Grading:** To get full credits solve any two of the questions.

### Question 1: Parallel Scaling (30 points)


A program simulates  $N$  particles, all particles interact with each other and, thus, the number of interactions is proportional to  $N^2$ . The runtimes of the program in seconds on  $P$  processor cores are in the table:

P \ N	500	1000	1500	2000
1	6.00	30.00	72.00	120.00
4	1.50	7.50	18.00	30.00
9	0.75	3.50	9.00	20.00
16	0.50	2.15	6.00	12.00
24	0.40	1.50	4.50	10.00

Plot four points of the strong scaling. Show all steps of the calculations. Label the axes.



Plot four points of the weak scaling using  $N = 500$  and  $P = 1$  as a reference. Show all steps of the calculations. Label the axes.



## Question 2: Diffusion (30 points)

The diffusion of a substance can be described by the equation

$$\frac{\partial c(x, y, t)}{\partial t} = D \left( \frac{\partial^2 c(x, y, t)}{\partial x^2} + \frac{\partial^2 c(x, y, t)}{\partial y^2} \right),$$

where  $c$  is the concentration of the substance at position  $(x, y)$  and at time  $t$ , and  $D$  is the diffusion constant. The diffusion process happens in the domain  $|x| < L/2$  and  $|y| < L/2$ . The concentration is zero on the boundaries of the domain. The initial concentration is

$$c(x, y, 0) = \begin{cases} 1, & \text{if } |x| < L/4 \text{ and } |y| < L/4, \\ 0, & \text{otherwise.} \end{cases}$$

- a) The skeleton code solves the equation on a uniform grid using a central finite difference scheme in space and forward Euler time integration. Parallellize the code by filling parts marked by `TODO` in the functions `advance` and `main`. Use a tiling decomposition scheme (i.e., distribute the rows evenly to the MPI processes).
- b) For a given time compute the integral of  $c(x, y, t)$  over the domain (total ammount of the substance). Fill the missing MPI parts in `compute_diagnostics`, and plot the result as a function of time using  $D = 1$ ,  $L = 2$  and  $N = 100$ .
- c) For a given time compute the histogram of  $c(x, y, t)$  in the function `compute_histogram` by implementing the missing MPI parts marked by `TODO`, and plot the resulting histogram for  $t = 0.5$  using  $D = 1$ ,  $L = 2$  and  $N = 100$ .




Figure 1: Left: Wake of an airplane visualized by condensation behind the engines. The vorticity sheet is generated at the trailing edge of the wings. Right: Vortex sheet at  $t = 1$  from the simulation.

### Question 3: Roll-up of a vortex line (30 points)

We want to simulate the evolution of a two-dimensional vorticity sheet similar to the wake of an airplane (Figure 1). We use  $N$  particles moving with time  $t$ , each particle  $i$  is located at  $(x_i(t), y_i(t))$  and carries a constant value  $\Gamma_i$  (circulation). The velocity field defined as

$$u(x, y, t) = \sum_{i=0}^{N-1} \frac{\Gamma_i}{2\pi} \frac{-[y - y_i(t)]}{\epsilon^2 + [x - x_i(t)]^2 + [y - y_i(t)]^2},$$

$$v(x, y, t) = \sum_{i=0}^{N-1} \frac{\Gamma_i}{2\pi} \frac{[x - x_i(t)]}{\epsilon^2 + [x - x_i(t)]^2 + [y - y_i(t)]^2}$$

is an approximate solution of Euler equations and corresponds to the vorticity field

$$\omega = \frac{\partial v(x, y, t)}{\partial x} - \frac{\partial u(x, y, t)}{\partial y} = \sum_{i=0}^{N-1} \Gamma_i \delta_\epsilon(x - x_i(t), y - y_i(t))$$

where  $\delta_\epsilon(x, y) = \frac{1}{\pi} \frac{\epsilon^2}{(\epsilon^2 + x^2 + y^2)^2}$  approximate a Dirac-delta function for a small  $\epsilon$ . A particle moves with the velocity at its location

$$\frac{dx_i(t)}{dt} = u(x_i(t), y_i(t), t), \quad \frac{dy_i(t)}{dt} = v(x_i(t), y_i(t), t),$$

Initially, particles are placed at  $y_i = 0$  and

$$x_i = -\frac{1}{2} + \frac{i + 1/2}{N}, \quad i = 0, \dots, N - 1$$

and have circulation  $\Gamma_i = \frac{1}{N} \frac{4x_i}{\sqrt{1-4x_i^2}}$

- Implement the interaction function `computeVelocities` in `q3/serial.cpp`. Check your results by visualizing the csv files with `paraview`.

- b) Parallelize your code using MPI by filling in the TODOs in `q3/mpi.cpp`. Each MPI rank must contain an equal number of particles. The parallelization of the `computeVelocities` can be done using a multi-pass communication, as described in the table:

		process $p$				
		0	1	$\dots$	P-2	P-1
	0	$D_0$	$D_1$	$\dots$	$D_{P-2}$	$D_{P-1}$
	1	$D_{P-1}$	$D_0$	$\dots$	$D_{P-3}$	$D_{P-2}$
pass $q$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$
	P-1	$D_1$	$D_2$	$\dots$	$D_{P-1}$	$D_0$

the data necessary to compute the velocities ( $x_i$ ,  $y_i$  and  $\Gamma_i$ ) must be communicated to every rank in a cyclic manner until every rank has computed the interactions between its own particles with every particles in the simulation. Parallelize the function `dumpToCsv` by gathering the data on the root.

- c) Use non blocking MPI routines to overlap the communication with computation. Write your solution in `q3/mpi_non_blocking.cpp`.

## Question 4: Roofline Model (30 points)

Given the following code:

```
1  float A[N], B[N], C[N];
2  ...
3  const int P = 2;
4  for (int i = 0; i < N; ++i) {
5      int j = 0;
6      while (j < P) {
7          A[i] = B[i] * A[i] + 0.5;
8          ++j;
9      }
10     C[i] = 0.9 * A[i] + C[i];
11 }
```

- What is the floating point operational intensity of the code? State all the assumptions you made and show your calculations.
- For a peak performance of 409.7 GFLOP/s (single precision) and a memory bandwidth of 34 GB/s, find all positive  $P$  for which the code is memory bound. Assume an infinite cache and state further assumptions you made. Show your calculations.
- Draw below the roofline corresponding to (b) and label the axes.

